

154HW4

Jiyeon Clover Jeong

10/11/2017

```
library(ElemStatLearn)
library(glmnet)

## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-13

library(leaps)
library(pls)

##
## Attaching package: 'pls'
## The following object is masked from 'package:stats':
##
##      loadings

library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
##
## Attaching package: 'caret'
## The following object is masked from 'package:pls':
##
##      R2
```

Problem 1 (10 pts)

$$r_{1,4} = \frac{\text{cov}(X1, X4)}{SD(X1)SD(X4)}$$

and $\text{cov}(X1, X4) = \text{cov}(X1, X1 + X2 + X3) = \text{cov}(X1, X1) + \text{cov}(X1, X2) + \text{cov}(X1, X3) = \text{var}(X1) + 0 + 0 = \text{var}(X1) = \sigma_1^2$

$\text{var}(X4) = \text{var}(X1 + X2 + X3) = \text{var}(X1) + \text{var}(X2) + \text{var}(X3) = \sigma_1^2 + \sigma_1^2 + \sigma_1^2 = 3\sigma_1^2$

Therefore,

$$r_{1,4} = \frac{\sigma_1^2}{\sqrt{\sigma_1^2 * 3\sigma_1^2}} = 1/\sqrt{3} = 0.5773503$$

$$r_{2,4} = \frac{\text{cov}(X2, X4)}{SD(X2)SD(X4)}$$

and $\text{cov}(X2, X4) = \text{cov}(X2, X1 + X2 + X3) = \text{cov}(X2, X1) + \text{cov}(X2, X2) + \text{cov}(X2, X3) = \text{var}(X2) = \sigma_1^2$

$\text{var}(X4) = \text{var}(X1 + X2 + X3) = \text{var}(X1) + \text{var}(X2) + \text{var}(X3) = \sigma_1^2 + \sigma_1^2 + \sigma_1^2 = 3\sigma_1^2$

Therefore,

$$r_{2,4} = \frac{\sigma_1^2}{\sqrt{\sigma_1^2 * 3\sigma_1^2}} = 1/\sqrt{3} = 0.5773503$$

$$r_{3,4} = \frac{cov(X3, X4)}{SD(X3)SD(X4)}$$

and $cov(X3, X4) = cov(X3, X1 + X2 + X3) = cov(X3, X1) + cov(X3, X2) + cov(X3, X3) = var(X3) = \sigma_1^2$
 $var(X4) = var(X1 + X2 + X3) = var(X1) + var(X2) + var(X3) = \sigma_1^2 + \sigma_1^2 + \sigma_1^2 = 3\sigma_1^2$

Therefore,

$$r_{3,4} = \frac{\sigma_1^2}{\sqrt{\sigma_1^2 * 3\sigma_1^2}} = 1/\sqrt{3} = 0.5773503$$

Problem 2

Show that any two components

$$z_h^T z_l = 0$$

for

$$(h \neq l)$$

are indeed orthogonal.

**** Base Case (i+1) ****

$$z_i^T z_{i+1} = z_i^T \left(\frac{X_i w_{i+1}}{w_{i+1}^T w_{i+1}} \right) = \frac{1}{w_{i+1}^T w_{i+1}} z_i^T (X_i w_{i+1}). \longrightarrow z_i^T (X_i w_{i+1}) = 0.$$

$$\text{And, } z_i^T (X_i w_{i+1}) = z_i^T ([x_{i-1} - z_i p_i^T] w_{i+1}) = z_i^T ([x_{i-1} - z_i [\frac{x_{i-1}^T z_i}{z_i^T z_i}]^T] w_{i+1}) = (z_i^T x_{i-1} - z_i^T x_{i-1}) w_{i+1} = 0.$$

**** Recursion (i+2) ****

$$z_i^T z_{i+2} = z_i^T (X_{i+1} w_{i+2}) \frac{1}{w_{i+2}^T w_{i+2}} = z_i^T (X_i - z_{i+1} p_{i+1}^T) \frac{w_{i+2}}{w_{i+2}^T w_{i+2}} = (z_i^T X_i - z_i^T z_{i+1} p_{i+1}^T) \frac{w_{i+2}}{w_{i+2}^T w_{i+2}}.$$

$$z_i^T z_{i+1} = 0 \text{ as we proved in the Base case.}$$

$$(z_i^T X_i - z_i^T z_{i+1} p_{i+1}^T) \frac{w_{i+2}}{w_{i+2}^T w_{i+2}} = z_i^T X_i \frac{w_{i+2}}{w_{i+2}^T w_{i+2}}.$$

$$\longrightarrow \text{Show } z_i^T X_i = 0.$$

$$z_i^T X_i = z_i^T (X_{i-1} - z_i p_i^T) = z_i^T (X_{i-1} - z_i [\frac{x_{i-1}^T z_i}{z_i^T z_i}]^T) = z_i^T X_{i-1} - z_i^T X_{i-1} = 0.$$

$$\longrightarrow z_i^T z_{i+2} = 0$$

Proof finished.

Problem 3

```
length(names(prostate))
```

```
## [1] 10
```

```
names(prostate)

## [1] "lcavol" "lweight" "age"      "lbph"      "svi"      "lcp"      "gleason"
## [8] "pgg45"  "lpsa"      "train"

train <- prostate[prostate$train == "TRUE", -10 ]

test <- prostate[prostate$train == "FALSE", -10 ]
```

Correlations of predictors, and some preprocessing (10 pts)

```
cor(train[, -9])

##          lcavol    lweight      age      lbph      svi      lcp
## lcavol  1.00000000  0.30023199  0.2863243  0.06316772  0.5929491  0.69204308
## lweight  0.30023199  1.00000000  0.3167235  0.43704154  0.1810545  0.15682859
## age      0.28632427  0.31672347  1.0000000  0.28734645  0.1289023  0.17295140
## lbph     0.06316772  0.43704154  0.2873464  1.00000000 -0.1391468 -0.08853456
## svi      0.59294913  0.18105448  0.1289023 -0.13914680  1.0000000  0.67124021
## lcp      0.69204308  0.15682859  0.1729514 -0.08853456  0.6712402  1.00000000
## gleason  0.42641407  0.02355821  0.3659151  0.03299215  0.3068754  0.47643684
## pgg45    0.48316136  0.07416632  0.2758057 -0.03040382  0.4813577  0.66253335
##          gleason      pgg45
## lcavol  0.42641407  0.48316136
## lweight  0.02355821  0.07416632
## age      0.36591512  0.27580573
## lbph     0.03299215 -0.03040382
## svi      0.30687537  0.48135774
## lcp      0.47643684  0.66253335
## gleason  1.00000000  0.75705650
## pgg45    0.75705650  1.00000000

train_stan <- scale(train[, -9])

summary(train_stan)

##          lcavol          lweight          age
## Min.      :-2.1411  Min.      :-2.62526  Min.      :-3.16524
## 1st Qu.   :-0.6641  1st Qu.   :-0.62054  1st Qu.   :-0.49935
## Median    : 0.1242  Median    :-0.05755  Median    : 0.03382
## Mean      : 0.0000  Mean      : 0.00000  Mean      : 0.00000
## 3rd Qu.   : 0.8334  3rd Qu.   : 0.54029  3rd Qu.   : 0.56700
## Max.      : 2.0180  Max.      : 2.42189  Max.      : 1.89994
##          lbph          svi          lcp          gleason
## Min.      :-0.99595  Min.      :-0.5331  Min.      :-0.8368  Min.      :-1.032
## 1st Qu.   :-0.99595  1st Qu.   :-0.5331  1st Qu.   :-0.8368  1st Qu.   :-1.032
## Median    :-0.08385  Median    :-0.5331  Median    :-0.4171  Median    : 0.379
## Mean      : 0.00000  Mean      : 0.0000  Mean      : 0.0000  Mean      : 0.000
## 3rd Qu.   : 1.00848  3rd Qu.   :-0.5331  3rd Qu.   : 0.8631  3rd Qu.   : 0.379
## Max.      : 1.54057  Max.      : 1.8480  Max.      : 2.0496  Max.      : 3.200
##          pgg45
## Min.      :-0.8965
## 1st Qu.   :-0.8965
## Median    :-0.3846
```

```
## Mean    : 0.0000
## 3rd Qu.: 0.8099
## Max.    : 2.5163
```

Least Squares Model (10 pts)

```
train_stan <- as.data.frame(cbind(train_stan, lpsa = train$lpsa))

lsfit <- lm(lpsa ~., data = train_stan)
summary(lsfit)$coefficients[,1:3]
```

```
##              Estimate Std. Error    t value
## (Intercept)  2.45234509 0.08701959 28.1815274
## lcavol       0.71640701 0.13350135  5.3662905
## lweight      0.29264240 0.10638488  2.7507894
## age         -0.14254963 0.10211957 -1.3959090
## lbph         0.21200760 0.10312428  2.0558456
## svi          0.30961953 0.12538985  2.4692552
## lcp         -0.28900562 0.15480404 -1.8669126
## gleason     -0.02091352 0.14257805 -0.1466812
## pgg45        0.27734595 0.15959237  1.7378397
```

Best Subset Regression (10 pts)

```
subset <- regsubsets(lpsa ~., data = train_stan)
coef(subset, 1:3)
```

```
## [[1]]
## (Intercept)      lcavol
##   2.4523451    0.8855136
##
## [[2]]
## (Intercept)      lcavol      lweight
##   2.4523451    0.7798589    0.3519101
##
## [[3]]
## (Intercept)      lcavol      lweight      svi
##   2.4523451    0.6461297    0.3511573    0.2259134
```

```
summary(subset)
```

```
## Subset selection object
## Call: regsubsets.formula(lpsa ~ ., data = train_stan)
## 8 Variables (and intercept)
##              Forced in Forced out
## lcavol      FALSE      FALSE
## lweight     FALSE      FALSE
## age         FALSE      FALSE
## lbph        FALSE      FALSE
## svi         FALSE      FALSE
## lcp         FALSE      FALSE
## gleason     FALSE      FALSE
```

```
## pgg45      FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##          lcavol lweight age lbph svi lcp gleason pgg45
## 1 ( 1 ) "*"      " "      " " " " " " " " " " " "
## 2 ( 1 ) "*"      "*"      " " " " " " " " " " " "
## 3 ( 1 ) "*"      "*"      " " " " "*" " " " " " " "
## 4 ( 1 ) "*"      "*"      " " "*" "*" " " " " " "
## 5 ( 1 ) "*"      "*"      " " "*" "*" " " " " "*"
## 6 ( 1 ) "*"      "*"      " " "*" "*" "*" " " " "*"
## 7 ( 1 ) "*"      "*"      "*" "*" "*" "*" " " " "*"
## 8 ( 1 ) "*"      "*"      "*" "*" "*" "*" "*" " "*"


```

```
# Best number of subset
```

```
minsubset <- which.min(summary(subset)$bic)
cat("The best number of subset is ", minsubset, "\n")
```

```
## The best number of subset is 2
```

```
plot(summary(subset)$bic, type = "o")
```



PCR and PLSR (40 pts)

```
set.seed(10)
```

```
pcr_model <- pcr(lpsa~., data = train_stan, validation = "CV")
```

```
# how to change number of folds in validation?
```

```
which.min(pcr_model$validation$PRESS)
```

```
## [1] 8
```

```
summary(pcr_model)
```

```
## Data:      X dimension: 67 8
## Y dimension: 67 1
## Fit method: svdpc
## Number of components considered: 8
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           1.217   0.9172   0.8975   0.8372   0.8323   0.8485   0.8467
## adjCV         1.217   0.9141   0.8949   0.8330   0.8269   0.8431   0.8419
##      7 comps  8 comps
## CV           0.8066   0.7833
## adjCV        0.8012   0.7767
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          42.83   63.24   76.20   83.92   89.61   94.32   97.82
## lpsa        45.18   50.84   59.58   61.00   61.17   62.08   66.36
##      8 comps
## X          100.00
## lpsa        69.44
```

```
# Model fits with the smallest CV-MSE of
```

```
cat("PCR Tuning parameter : ", which.min(pcr_model$validation$PRESS))
```

```
## PCR Tuning parameter : 8
```

```
print("Associated coefficients : ")
```

```
## [1] "Associated coefficients : "
```

```
pcr_model$coefficients[, , which.min(pcr_model$validation$PRESS)]
```

```
##      lcavol      lweight      age      lbph      svi      lcp
## 0.71640701 0.29264240 -0.14254963 0.21200760 0.30961953 -0.28900562
##      gleason      pgg45
## -0.02091352 0.27734595
```

```
plsr_model <- plsr(lpsa~., data = train_stan, validation = "CV")
summary(plsr_model)
```

```
## Data:      X dimension: 67 8
## Y dimension: 67 1
## Fit method: kernelpls
## Number of components considered: 8
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           1.217   0.8475   0.7885   0.7709   0.7616   0.7530   0.7524
## adjCV         1.217   0.8453   0.7844   0.7667   0.7559   0.7482   0.7476
```

```

##          7 comps  8 comps
## CV          0.7520  0.7521
## adjCV       0.7473  0.7474
##
## TRAINING: % variance explained
##          1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           41.64   58.29   71.13   79.75   86.08   90.21   94.70
## lpsa        55.79   64.60   67.51   69.12   69.37   69.43   69.44
##          8 comps
## X           100.00
## lpsa        69.44

plsr_model

## Partial least squares regression , fitted with the kernel algorithm.
## Cross-validated using 10 random segments.
## Call:
## plsr(formula = lpsa ~ ., data = train_stan, validation = "CV")
# Model fits with the smallest CV-MSE of

# Q how to get CV table from summary???

# CV MSE : MSEP(pcr_model)

cat("Plsr Tuning parameter : ", which.min(plsr_model$validation$PRESS))

## Plsr Tuning parameter : 7

print("Associated coefficients : ")

## [1] "Associated coefficients : "

plsr_model$coefficients[, , which.min(plsr_model$validation$PRESS)]

##          lcavol      lweight          age          lbph          svi          lcp
## 0.71636186 0.29192059 -0.14233890 0.21259719 0.31026935 -0.28905136
##          gleason      pgg45
## -0.02101011 0.27706347

# fit = glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa)
# plot(fit)

# Standardized coefficients??


# Q is this plot correct --- >standardized coefficients ???

# coefplot(pcr_model, ncomp = 1:8, legendpos = "bottomright", xlab = "number of components")
# plot(pcr_model, plottype = "coefficients", ncomp = 1:8)


pcr_coefs = apply(pcr_model$coefficients, 3, function(x) x)

matplot(t(pcr_coefs), type= 'l', lwd = 2, xlab = "Number of Components", ylab = "Standardized Coefficients",
legend("topleft", colnames(pcr_coefs),col=seq_len(ncol(pcr_coefs)),cex=0.4,fill=seq_len(ncol(pcr_coefs)))

```




```
matplot(t(pcr_coefs), type= 's', lwd = 2, xlab = "Number of Components", ylab = "Standardized Coefficients",
legend("topleft", colnames(pcr_coefs),col=seq_len(ncol(pcr_coefs)),cex=0.4,fill=seq_len(ncol(pcr_coefs)))
```



```
validationplot(pcr_model, val.type = "MSEP", ylab = "CV-MSE")
```


Ipsa




```
#plot(pcr_model$validation$PRESS[1,], type = "o", main = "PCR")
```

```
# coefplot(plsr_model, ncomp = 1:8, legendpos = "bottomright")
```


```
# plot(plsr_model, plottype = "coefficients", ncomp = 1:8)
```

```
plsr_coefs = apply(plsr_model$coefficients, 3, function(x) x)
```


```
matplot(t(plsr_coefs), type= 'l', lwd = 2, xlab = "Number of Components", ylab = "Standardized Coefficients",
legend("topleft", colnames(plsr_coefs),col=seq_len(ncol(plsr_coefs)),cex=0.4,fill=seq_len(ncol(plsr_coefs)))
```



```
matplot(t(plsr_coefs), type= 's', lwd = 2, xlab = "Number of Components", ylab = "Standardized Coefficients",
legend("topleft", colnames(plsr_coefs),col=seq_len(ncol(plsr_coefs)),cex=0.4,fill=seq_len(ncol(plsr_coefs)))
```



```
validationplot(plsr_model, val.type = "MSEP", ylab = "CV-MSE")
```



RR and Lasso (40 pts)

```
set.seed(10)

# Fitting the model (Ridge: Alpha = 0)
ridgecv <- cv.glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 0)

ridgecv
```

```
## $lambda
## [1] 878.88041366 800.80310001 729.66195971 664.84080224 605.77817775
## [6] 551.96251403 502.92768556 458.24897610 417.53940005 380.44635054
## [11] 346.64854531 315.85324395 287.79371228 262.22691208 238.93139593
## [16] 217.70538924 198.36504248 180.74283883 164.68614317 150.05588011
## [21] 136.72532929 124.57902787 113.51177039 103.42769756 94.23946596
## [26] 85.86749154 78.23926025 71.28870001 64.95560838 59.18513116
## [31] 53.92728723 49.13653566 44.77138126 40.79401514 37.16998727
## [36] 33.86790804 30.85917644 28.11773225 25.61983042 23.34383529
## [41] 21.27003330 19.38046217 17.65875532 16.09000017 14.66060890
## [46] 13.35820081 12.17149507 11.09021299 10.10498904 9.20728966
## [51] 8.38933942 7.64405363 6.96497698 6.34622762 5.78244625
## [56] 5.26874966 4.80068846 4.37420853 3.98561590 3.63154477
## [61] 3.30892834 3.01497227 2.74713044 2.50308294 2.28071595
## [66] 2.07810343 1.89349045 1.72527798 1.57200904 1.43235609
## [71] 1.30510952 1.18916719 1.08352485 0.98726749 0.89956137
## [76] 0.81964682 0.74683165 0.68048519 0.62003276 0.56495076
## [81] 0.51476209 0.46903204 0.42736453 0.38939864 0.35480554
## [86] 0.32328559 0.29456579 0.26839738 0.24455370 0.22282822
```

```

## [91] 0.20303277 0.18499590 0.16856138 0.15358685 0.13994262
## [96] 0.12751050 0.11618282 0.10586146 0.09645702
##
## $cvm
## [1] 1.4624171 1.4588464 1.4557350 1.4547654 1.4539582 1.4530739 1.4521052
## [8] 1.4510443 1.4498826 1.4486110 1.4472192 1.4456963 1.4440305 1.4422090
## [15] 1.4402178 1.4380421 1.4356658 1.4330714 1.4302405 1.4271531 1.4237880
## [22] 1.4201226 1.4161330 1.4117939 1.4070786 1.4019592 1.3964066 1.3903908
## [29] 1.3838807 1.3768449 1.3692513 1.3610681 1.3522638 1.3428079 1.3326713
## [36] 1.3218273 1.3102518 1.2979243 1.2848292 1.2709557 1.2562998 1.2408642
## [43] 1.2246599 1.2077064 1.1900326 1.1716772 1.1526892 1.1331273 1.1130601
## [50] 1.0925653 1.0717288 1.0506434 1.0294073 1.0081221 0.9868914 0.9658180
## [57] 0.9450023 0.9245401 0.9045207 0.8850256 0.8661268 0.8478860 0.8303540
## [64] 0.8135704 0.7975639 0.7823526 0.7679446 0.7543392 0.7415273 0.7294934
## [71] 0.7182160 0.7076692 0.6978235 0.6886473 0.6801064 0.6721697 0.6647990
## [78] 0.6579699 0.6516401 0.6457817 0.6403657 0.6353647 0.6307637 0.6265293
## [85] 0.6226447 0.6190893 0.6158408 0.6128857 0.6102119 0.6078035 0.6056559
## [92] 0.6037363 0.6020391 0.6005525 0.5992583 0.5981385 0.5971935 0.5964014
## [99] 0.5957481
##
## $cvstd
## [1] 0.25120463 0.25152114 0.24997254 0.24970794 0.24957030 0.24941941
## [7] 0.24925403 0.24907278 0.24887416 0.24865655 0.24841816 0.24815706
## [13] 0.24787115 0.24755812 0.24721549 0.24684055 0.24643039 0.24598182
## [19] 0.24549140 0.24495545 0.24436995 0.24373062 0.24303284 0.24227166
## [25] 0.24144179 0.24053760 0.23955310 0.23848194 0.23731744 0.23605255
## [31] 0.23467994 0.23319195 0.23158069 0.22983804 0.22795578 0.22592557
## [37] 0.22373914 0.22138835 0.21886533 0.21616264 0.21327343 0.21019165
## [43] 0.20691223 0.20343136 0.19974667 0.19585749 0.19176518 0.18747332
## [49] 0.18298797 0.17831793 0.17347494 0.16847383 0.16333271 0.15807299
## [55] 0.15271947 0.14730021 0.14184646 0.13639241 0.13097493 0.12563308
## [61] 0.12040765 0.11534052 0.11047385 0.10584916 0.10150632 0.09748226
## [67] 0.09380970 0.09051572 0.08762043 0.08513578 0.08306471 0.08140088
## [73] 0.08012883 0.07922498 0.07865857 0.07839414 0.07839664 0.07862499
## [79] 0.07904075 0.07960777 0.08029215 0.08106522 0.08190291 0.08277949
## [85] 0.08367818 0.08458392 0.08548206 0.08636460 0.08722548 0.08805939
## [91] 0.08887153 0.08964818 0.09038762 0.09109286 0.09176492 0.09240508
## [97] 0.09301170 0.09358723 0.09413297
##
## $cvup
## [1] 1.7136218 1.7103675 1.7057076 1.7044733 1.7035285 1.7024933 1.7013592
## [8] 1.7001171 1.6987568 1.6972675 1.6956373 1.6938534 1.6919017 1.6897671
## [15] 1.6874333 1.6848827 1.6820962 1.6790532 1.6757319 1.6721085 1.6681579
## [22] 1.6638532 1.6591659 1.6540656 1.6485204 1.6424968 1.6359597 1.6288727
## [29] 1.6211982 1.6128974 1.6039312 1.5942600 1.5838445 1.5726459 1.5606271
## [36] 1.5477529 1.5339909 1.5193127 1.5036945 1.4871184 1.4695732 1.4510559
## [43] 1.4315722 1.4111377 1.3897793 1.3675347 1.3444544 1.3206006 1.2960481
## [50] 1.2708832 1.2452038 1.2191172 1.1927400 1.1661951 1.1396109 1.1131182
## [57] 1.0868487 1.0609325 1.0354957 1.0106587 0.9865345 0.9632265 0.9408278
## [64] 0.9194195 0.8990702 0.8798349 0.8617543 0.8448549 0.8291477 0.8146292
## [71] 0.8012807 0.7890700 0.7779524 0.7678723 0.7587650 0.7505639 0.7431956
## [78] 0.7365949 0.7306809 0.7253895 0.7206578 0.7164300 0.7126666 0.7093088
## [85] 0.7063229 0.7036732 0.7013228 0.6992503 0.6974374 0.6958629 0.6945275
## [92] 0.6933845 0.6924268 0.6916454 0.6910232 0.6905436 0.6902052 0.6899887

```

```

## [99] 0.6898811
##
## $cvlo
## [1] 1.2112125 1.2073252 1.2057625 1.2050575 1.2043879 1.2036545 1.2028512
## [8] 1.2019715 1.2010085 1.1999544 1.1988010 1.1975392 1.1961594 1.1946509
## [15] 1.1930024 1.1912016 1.1892354 1.1870896 1.1847491 1.1821976 1.1794180
## [22] 1.1763920 1.1731002 1.1695222 1.1656368 1.1614216 1.1568535 1.1519088
## [29] 1.1465633 1.1407923 1.1345713 1.1278761 1.1206831 1.1129698 1.1047156
## [36] 1.0959017 1.0865126 1.0765360 1.0659638 1.0547931 1.0430263 1.0306726
## [43] 1.0177477 1.0042750 0.9902859 0.9758198 0.9609240 0.9456540 0.9300721
## [50] 0.9142474 0.8982539 0.8821696 0.8660746 0.8500491 0.8341719 0.8185178
## [57] 0.8031558 0.7881476 0.7735458 0.7593926 0.7457192 0.7325455 0.7198801
## [64] 0.7077212 0.6960576 0.6848703 0.6741349 0.6638234 0.6539069 0.6443576
## [71] 0.6351513 0.6262683 0.6176947 0.6094223 0.6014479 0.5937756 0.5864024
## [78] 0.5793449 0.5725994 0.5661740 0.5600735 0.5542995 0.5488608 0.5437498
## [85] 0.5389665 0.5345053 0.5303587 0.5265211 0.5229865 0.5197442 0.5167844
## [92] 0.5140882 0.5116515 0.5094596 0.5074933 0.5057334 0.5041818 0.5028142
## [99] 0.5016152
##
## $nzero
## s0 s1 s2 s3 s4 s5 s6 s7 s8 s9 s10 s11 s12 s13 s14 s15 s16 s17
## 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
## s18 s19 s20 s21 s22 s23 s24 s25 s26 s27 s28 s29 s30 s31 s32 s33 s34 s35
## 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
## s36 s37 s38 s39 s40 s41 s42 s43 s44 s45 s46 s47 s48 s49 s50 s51 s52 s53
## 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
## s54 s55 s56 s57 s58 s59 s60 s61 s62 s63 s64 s65 s66 s67 s68 s69 s70 s71
## 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
## s72 s73 s74 s75 s76 s77 s78 s79 s80 s81 s82 s83 s84 s85 s86 s87 s88 s89
## 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
## s90 s91 s92 s93 s94 s95 s96 s97 s98
## 8 8 8 8 8 8 8 8 8
##
## $name
## mse
## "Mean-Squared Error"
##
## $glmnet.fit
##
## Call: glmnet(x = as.matrix(train_stan[, -9]), y = train_stan$lpsa, alpha = 0)
##
## Df %Dev Lambda
## [1,] 8 3.558e-36 878.90000
## [2,] 8 5.236e-03 800.80000
## [3,] 8 5.743e-03 729.70000
## [4,] 8 6.298e-03 664.80000
## [5,] 8 6.906e-03 605.80000
## [6,] 8 7.573e-03 552.00000
## [7,] 8 8.303e-03 502.90000
## [8,] 8 9.102e-03 458.20000
## [9,] 8 9.978e-03 417.50000
## [10,] 8 1.094e-02 380.40000
## [11,] 8 1.199e-02 346.60000
## [12,] 8 1.313e-02 315.90000

```

```

## [13,] 8 1.439e-02 287.80000
## [14,] 8 1.576e-02 262.20000
## [15,] 8 1.726e-02 238.90000
## [16,] 8 1.890e-02 217.70000
## [17,] 8 2.070e-02 198.40000
## [18,] 8 2.265e-02 180.70000
## [19,] 8 2.479e-02 164.70000
## [20,] 8 2.711e-02 150.10000
## [21,] 8 2.965e-02 136.70000
## [22,] 8 3.242e-02 124.60000
## [23,] 8 3.543e-02 113.50000
## [24,] 8 3.870e-02 103.40000
## [25,] 8 4.226e-02 94.24000
## [26,] 8 4.612e-02 85.87000
## [27,] 8 5.031e-02 78.24000
## [28,] 8 5.485e-02 71.29000
## [29,] 8 5.976e-02 64.96000
## [30,] 8 6.507e-02 59.19000
## [31,] 8 7.081e-02 53.93000
## [32,] 8 7.699e-02 49.14000
## [33,] 8 8.364e-02 44.77000
## [34,] 8 9.078e-02 40.79000
## [35,] 8 9.845e-02 37.17000
## [36,] 8 1.066e-01 33.87000
## [37,] 8 1.154e-01 30.86000
## [38,] 8 1.247e-01 28.12000
## [39,] 8 1.346e-01 25.62000
## [40,] 8 1.451e-01 23.34000
## [41,] 8 1.562e-01 21.27000
## [42,] 8 1.679e-01 19.38000
## [43,] 8 1.802e-01 17.66000
## [44,] 8 1.931e-01 16.09000
## [45,] 8 2.065e-01 14.66000
## [46,] 8 2.205e-01 13.36000
## [47,] 8 2.349e-01 12.17000
## [48,] 8 2.498e-01 11.09000
## [49,] 8 2.651e-01 10.10000
## [50,] 8 2.808e-01 9.20700
## [51,] 8 2.967e-01 8.38900
## [52,] 8 3.128e-01 7.64400
## [53,] 8 3.291e-01 6.96500
## [54,] 8 3.455e-01 6.34600
## [55,] 8 3.618e-01 5.78200
## [56,] 8 3.781e-01 5.26900
## [57,] 8 3.942e-01 4.80100
## [58,] 8 4.100e-01 4.37400
## [59,] 8 4.256e-01 3.98600
## [60,] 8 4.408e-01 3.63200
## [61,] 8 4.556e-01 3.30900
## [62,] 8 4.699e-01 3.01500
## [63,] 8 4.838e-01 2.74700
## [64,] 8 4.971e-01 2.50300
## [65,] 8 5.099e-01 2.28100
## [66,] 8 5.221e-01 2.07800

```

```
## [67,] 8 5.337e-01 1.89300
## [68,] 8 5.447e-01 1.72500
## [69,] 8 5.552e-01 1.57200
## [70,] 8 5.651e-01 1.43200
## [71,] 8 5.745e-01 1.30500
## [72,] 8 5.833e-01 1.18900
## [73,] 8 5.917e-01 1.08400
## [74,] 8 5.995e-01 0.98730
## [75,] 8 6.068e-01 0.89960
## [76,] 8 6.137e-01 0.81960
## [77,] 8 6.202e-01 0.74680
## [78,] 8 6.263e-01 0.68050
## [79,] 8 6.319e-01 0.62000
## [80,] 8 6.372e-01 0.56500
## [81,] 8 6.422e-01 0.51480
## [82,] 8 6.468e-01 0.46900
## [83,] 8 6.511e-01 0.42740
## [84,] 8 6.551e-01 0.38940
## [85,] 8 6.589e-01 0.35480
## [86,] 8 6.623e-01 0.32330
## [87,] 8 6.655e-01 0.29460
## [88,] 8 6.685e-01 0.26840
## [89,] 8 6.712e-01 0.24460
## [90,] 8 6.737e-01 0.22280
## [91,] 8 6.759e-01 0.20300
## [92,] 8 6.780e-01 0.18500
## [93,] 8 6.799e-01 0.16860
## [94,] 8 6.816e-01 0.15360
## [95,] 8 6.831e-01 0.13990
## [96,] 8 6.845e-01 0.12750
## [97,] 8 6.858e-01 0.11620
## [98,] 8 6.869e-01 0.10590
## [99,] 8 6.878e-01 0.09646
## [100,] 8 6.887e-01 0.08789
```

```
##
## $lambda.min
## [1] 0.09645702
##
## $lambda.1se
## [1] 0.9872675
##
## attr("class")
## [1] "cv.glmnet"
```

```
coef(ridgecv, s = "lambda.min")
```

```
## 9 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 2.45234509
## lcavol      0.60438317
## lweight     0.28576500
## age         -0.10858418
## lbph        0.20096586
## svi         0.28336365
## lcp         -0.15469409
```

```
## gleason      0.01414138
## pgg45        0.20305366
```

```
opt_lambda <- ridgecv$lambda.min
opt_lambda
```

```
## [1] 0.09645702
```

```
# Q lambda.min.ratio = opt_lambda
```

```
ridge <- glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 0, lambda.min.ratio = opt_lambda)
```

```
ridge <- glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 0, lambda = opt_lambda)
```

```
# Q is summary enough???
summary(ridge)
```

```
##          Length Class      Mode
## a0         1      -none-   numeric
## beta       8      dgCMatrx S4
## df         1      -none-   numeric
## dim        2      -none-   numeric
## lambda     1      -none-   numeric
## dev.ratio  1      -none-   numeric
## nulldev    1      -none-   numeric
## npasses    1      -none-   numeric
## jerr       1      -none-   numeric
## offset     1      -none-   logical
## call       5      -none-   call
## nobs       1      -none-   numeric
```


```
# Q : ridge.coef gives the best model's coefficients???
```

```
coef(ridgecv, s = ridgecv$lambda.min)
```

```
## 9 x 1 sparse Matrix of class "dgCMatrx"
```

```
##              1
## (Intercept)  2.45234509
## lcavol       0.60438317
## lweight      0.28576500
## age         -0.10858418
## lbph         0.20096586
## svi          0.28336365
## lcp          -0.15469409
## gleason      0.01414138
## pgg45        0.20305366
```


```
plot.cv.glmnet(ridgecv)
```

```
# Q :      ??? straight?
```

```
ridge <- glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 0)
```

```
plot.glmnet(ridge)
```



```

set.seed(10)
# Fitting the model (Lasso: Alpha = 1)

lassocv <- cv.glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 1)

lassocv

## $lambda
## [1] 0.878880414 0.800803100 0.729661960 0.664840802 0.605778178
## [6] 0.551962514 0.502927686 0.458248976 0.417539400 0.380446351
## [11] 0.346648545 0.315853244 0.287793712 0.262226912 0.238931396
## [16] 0.217705389 0.198365042 0.180742839 0.164686143 0.150055880
## [21] 0.136725329 0.124579028 0.113511770 0.103427698 0.094239466
## [26] 0.085867492 0.078239260 0.071288700 0.064955608 0.059185131
## [31] 0.053927287 0.049136536 0.044771381 0.040794015 0.037169987
## [36] 0.033867908 0.030859176 0.028117732 0.025619830 0.023343835
## [41] 0.021270033 0.019380462 0.017658755 0.016090000 0.014660609
## [46] 0.013358201 0.012171495 0.011090213 0.010104989 0.009207290
## [51] 0.008389339 0.007644054 0.006964977 0.006346228 0.005782446
## [56] 0.005268750 0.004800688 0.004374209 0.003985616 0.003631545
## [61] 0.003308928 0.003014972 0.002747130 0.002503083 0.002280716
## [66] 0.002078103 0.001893490 0.001725278
##
## $cvm
## [1] 1.4495864 1.3748827 1.2661084 1.1715141 1.0929790 1.0277767 0.9748502
## [8] 0.9308870 0.8914314 0.8559751 0.8194052 0.7864824 0.7597421 0.7381668
## [15] 0.7217217 0.7078173 0.6958802 0.6846987 0.6730528 0.6629171 0.6532195
## [22] 0.6450963 0.6382981 0.6312325 0.6252097 0.6201932 0.6165538 0.6152285
## [29] 0.6146022 0.6143005 0.6148388 0.6166678 0.6174565 0.6166245 0.6156366
## [36] 0.6148322 0.6136893 0.6108610 0.6082566 0.6061759 0.6043631 0.6023808
## [43] 0.6004946 0.5991178 0.5981771 0.5974701 0.5969291 0.5965204 0.5962029
## [50] 0.5959843 0.5958515 0.5957964 0.5958097 0.5958886 0.5959800 0.5960695
## [57] 0.5961649 0.5962521 0.5963552 0.5964501 0.5965507 0.5966405 0.5967238
## [64] 0.5968069 0.5968945 0.5969681 0.5970432 0.5971159
##
## $cvstd
## [1] 0.25327067 0.26159165 0.23977068 0.21797620 0.19887861 0.18207992
## [7] 0.16685341 0.15395216 0.14365053 0.13530767 0.12572314 0.11688194
## [13] 0.10931720 0.10291853 0.09721438 0.09218828 0.08813678 0.08510700
## [19] 0.08298359 0.08147538 0.08009646 0.07929736 0.07896562 0.07827778
## [25] 0.07802900 0.07818849 0.07872737 0.07964203 0.08081182 0.08208932
## [31] 0.08351785 0.08503489 0.08663478 0.08822121 0.08960356 0.09084344
## [37] 0.09172827 0.09222408 0.09257572 0.09296872 0.09334075 0.09355016
## [43] 0.09369661 0.09401804 0.09447775 0.09496337 0.09543752 0.09589600
## [49] 0.09633315 0.09675071 0.09714485 0.09751823 0.09786233 0.09817762
## [55] 0.09847383 0.09874705 0.09900509 0.09924459 0.09946556 0.09966971
## [61] 0.09986063 0.10003448 0.10019260 0.10034093 0.10047097 0.10059496
## [67] 0.10070746 0.10081329
##
## $cvup
## [1] 1.7028571 1.6364744 1.5058791 1.3894903 1.2918576 1.2098566 1.1417036
## [8] 1.0848391 1.0350820 0.9912827 0.9451284 0.9033644 0.8690593 0.8410853
## [15] 0.8189361 0.8000056 0.7840170 0.7698057 0.7560364 0.7443924 0.7333159

```

```

## [22] 0.7243936 0.7172637 0.7095103 0.7032387 0.6983817 0.6952812 0.6948705
## [29] 0.6954140 0.6963899 0.6983566 0.7017027 0.7040912 0.7048457 0.7052401
## [36] 0.7056756 0.7054176 0.7030851 0.7008324 0.6991446 0.6977039 0.6959310
## [43] 0.6941912 0.6931358 0.6926548 0.6924335 0.6923667 0.6924164 0.6925361
## [50] 0.6927350 0.6929963 0.6933146 0.6936720 0.6940662 0.6944538 0.6948165
## [57] 0.6951700 0.6954967 0.6958207 0.6961198 0.6964113 0.6966750 0.6969164
## [64] 0.6971478 0.6973654 0.6975631 0.6977507 0.6979292
##
## $cvlo
## [1] 1.1963158 1.1132911 1.0263377 0.9535379 0.8941004 0.8456968 0.8079968
## [8] 0.7769348 0.7477809 0.7206674 0.6936821 0.6696005 0.6504249 0.6352482
## [15] 0.6245073 0.6156291 0.6077434 0.5995917 0.5900692 0.5814417 0.5731230
## [22] 0.5657989 0.5593324 0.5529547 0.5471807 0.5420048 0.5378264 0.5355865
## [29] 0.5337903 0.5322112 0.5313209 0.5316329 0.5308217 0.5284033 0.5260330
## [36] 0.5239888 0.5219611 0.5186369 0.5156809 0.5132071 0.5110224 0.5088307
## [43] 0.5067979 0.5050997 0.5036993 0.5025068 0.5014916 0.5006244 0.4998698
## [50] 0.4992336 0.4987067 0.4982782 0.4979474 0.4977109 0.4975062 0.4973224
## [57] 0.4971599 0.4970075 0.4968896 0.4967804 0.4966901 0.4966060 0.4965311
## [64] 0.4964660 0.4964235 0.4963732 0.4963358 0.4963026
##
## $nzero
## s0 s1 s2 s3 s4 s5 s6 s7 s8 s9 s10 s11 s12 s13 s14 s15 s16 s17
## 0 1 1 1 1 1 1 1 2 2 3 3 3 3 3 3 5 5
## s18 s19 s20 s21 s22 s23 s24 s25 s26 s27 s28 s29 s30 s31 s32 s33 s34 s35
## 5 5 5 5 5 5 5 5 5 5 5 6 6 6 7 7 7 7
## s36 s37 s38 s39 s40 s41 s42 s43 s44 s45 s46 s47 s48 s49 s50 s51 s52 s53
## 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7
## s54 s55 s56 s57 s58 s59 s60 s61 s62 s63 s64 s65 s66 s67
## 7 7 8 8 8 8 8 8 8 8 8 8 8 8
##
## $name
## mse
## "Mean-Squared Error"
##
## $glmnet.fit
##
## Call: glmnet(x = as.matrix(train_stan[, -9]), y = train_stan$lpsa, alpha = 1)
##
## Df %Dev Lambda
## [1,] 0 0.00000 0.878900
## [2,] 1 0.09126 0.800800
## [3,] 1 0.16700 0.729700
## [4,] 1 0.22990 0.664800
## [5,] 1 0.28220 0.605800
## [6,] 1 0.32550 0.552000
## [7,] 1 0.36150 0.502900
## [8,] 1 0.39140 0.458200
## [9,] 2 0.42810 0.417500
## [10,] 2 0.45980 0.380400
## [11,] 3 0.48770 0.346600
## [12,] 3 0.51310 0.315900
## [13,] 3 0.53420 0.287800
## [14,] 3 0.55180 0.262200
## [15,] 3 0.56630 0.238900

```

```

## [16,] 3 0.57840 0.217700
## [17,] 5 0.59170 0.198400
## [18,] 5 0.60450 0.180700
## [19,] 5 0.61510 0.164700
## [20,] 5 0.62390 0.150100
## [21,] 5 0.63120 0.136700
## [22,] 5 0.63720 0.124600
## [23,] 5 0.64230 0.113500
## [24,] 5 0.64650 0.103400
## [25,] 5 0.64990 0.094240
## [26,] 5 0.65280 0.085870
## [27,] 5 0.65520 0.078240
## [28,] 5 0.65720 0.071290
## [29,] 5 0.65890 0.064960
## [30,] 6 0.66050 0.059190
## [31,] 6 0.66320 0.053930
## [32,] 6 0.66530 0.049140
## [33,] 7 0.66760 0.044770
## [34,] 7 0.67210 0.040790
## [35,] 7 0.67590 0.037170
## [36,] 7 0.67900 0.033870
## [37,] 7 0.68160 0.030860
## [38,] 7 0.68370 0.028120
## [39,] 7 0.68550 0.025620
## [40,] 7 0.68700 0.023340
## [41,] 7 0.68820 0.021270
## [42,] 7 0.68930 0.019380
## [43,] 7 0.69010 0.017660
## [44,] 7 0.69080 0.016090
## [45,] 7 0.69140 0.014660
## [46,] 7 0.69190 0.013360
## [47,] 7 0.69230 0.012170
## [48,] 7 0.69260 0.011090
## [49,] 7 0.69290 0.010100
## [50,] 7 0.69310 0.009207
## [51,] 7 0.69330 0.008389
## [52,] 7 0.69350 0.007644
## [53,] 7 0.69360 0.006965
## [54,] 7 0.69370 0.006346
## [55,] 7 0.69380 0.005782
## [56,] 7 0.69390 0.005269
## [57,] 8 0.69390 0.004801
## [58,] 8 0.69400 0.004374
## [59,] 8 0.69410 0.003986
## [60,] 8 0.69410 0.003632
## [61,] 8 0.69420 0.003309
## [62,] 8 0.69420 0.003015
## [63,] 8 0.69420 0.002747
## [64,] 8 0.69430 0.002503
## [65,] 8 0.69430 0.002281
## [66,] 8 0.69430 0.002078
## [67,] 8 0.69430 0.001893
## [68,] 8 0.69430 0.001725
## [69,] 8 0.69430 0.001572

```

```

## [70,] 8 0.69430 0.001432
## [71,] 8 0.69430 0.001305
##
## $lambda.min
## [1] 0.007644054
##
## $lambda.1se
## [1] 0.1807428
##
## attr("class")
## [1] "cv.glmnet"

opt_lambda <- lasso_cv$lambda.min
opt_lambda

## [1] 0.007644054

lasso <- glmnet(as.matrix(train_stan[, -9]), train_stan$lpsa , alpha = 1, lambda.min.ratio = opt_lambda)

lasso <- glmnet(as.matrix(train_stan[, -9]), train_stan$lpsa , alpha = 1, lambda = opt_lambda)

lasso

##
## Call:  glmnet(x = as.matrix(train_stan[, -9]), y = train_stan$lpsa,      alpha = 1, lambda = opt_lambda)
##
##      Df    %Dev   Lambda
## [1,]  7 0.6935 0.007644

coef(lasso_cv, s = "lambda.min")


## 9 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  2.4523451
## lcavol      0.6918697
## lweight     0.2887031
## age         -0.1268621
## lbph        0.2033674
## svi         0.2940763
## lcp         -0.2389979
## gleason     .
## pgg45       0.2357199

coef(lasso, s = "lambda.min")

## 9 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  2.4523451
## lcavol      0.6919179
## lweight     0.2887836
## age         -0.1269113
## lbph        0.2033150
## svi         0.2940639
## lcp         -0.2393184
## gleason     .
## pgg45       0.2359208


```

```
plot.cv.glmnet(lassocv)
```



```
lasso <- glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 1)
```

```
plot.glmnet(lasso)
```



Model Selection (20 pts)

```
subset.coef <- coef(subset, minsubset)

# c(coef(subset, minsubset), rep(NA,6))

ridge.coef <- (coef(ridgecv, s = ridgecv$lambda.1se))[1:9]
# Q : s = lambda.min?
# or coef(ridgecv)

lasso.coef <- (coef(lassocv, s = lasso$lambda.1se))[1:9]

# don't recycle
# Q PCR, PLS intercept coefficient needed

table <- cbind(LS = lsfit$coefficients,
               Best_Subset = c(subset.coef, rep(0, 6)),
               Ridge = ridge.coef , Lasso = lasso.coef, PCR = coef(pcr_model, intercept = T), PLS = coef(pls_model, intercept = T))

table
```

##		LS	Best_Subset		Ridge	Lasso	PCR
##	(Intercept)	2.45234509	2.4523451	2.452345085	2.45234509	2.45234509	2.45234509
##	lcavol	0.71640701	0.7798589	0.317607465	0.56536414	0.71640701	0.71640701
##	lweight	0.29264240	0.3519101	0.207749630	0.19936439	0.29264240	0.29264240
##	age	-0.14254963	0.0000000	-0.003245594	0.00000000	-0.14254963	-0.14254963
##	lbph	0.21200760	0.0000000	0.132737031	0.02917609	0.21200760	0.21200760
##	svi	0.30961953	0.0000000	0.192527220	0.11532140	0.30961953	0.30961953
##	lcp	-0.28900562	0.0000000	0.067982376	0.00000000	-0.28900562	-0.28900562
##	gleason	-0.02091352	0.0000000	0.052456497	0.00000000	-0.02091352	-0.02091352
##	pgg45	0.27734595	0.0000000	0.109563765	0.01647872	0.27734595	0.27734595

##		PLS
##	(Intercept)	2.45234509
##	lcavol	0.71640701
##	lweight	0.29264240
##	age	-0.14254963
##	lbph	0.21200760
##	svi	0.30961953
##	lcp	-0.28900562
##	gleason	-0.02091352
##	pgg45	0.27734595

```
y <- test[,9]
X <- as.data.frame(scale(test[,-9]))

LSmse <- mean((predict(lsfit, X) - y)^2)
# why so huge?
LSmse
```

```
## [1] 0.5491941
```

```

# Q : what's wrong??
#subsetmse <- mean(summary(subset)$rss^2)
#subset coefficients are the same as the OLS regression coefficients with the selected variable
subset_lm <- lm(lpsa ~ lcavol + lweight, data = train_stan)
subsetmse <- mean((predict(subset_lm, X) - y)^2)
subsetmse

## [1] 0.5483947

#yhat <- as.matrix(cbind(1, test[,c(1,2)])) %*% as.matrix(subset_coef)
# predict(subset_lm, test) give the same results so it should be correct.
#head(yhat)

#ridgecv <- cv.glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 0)
#opt_lambda <- ridgecv$lambda.min

# Q lambda.min.ratio = opt_lambda
#ridge <- glmnet(as.matrix(train_stan[,-9]), train_stan$lpsa , alpha = 0, lambda.min.ratio = opt_lambda)

#ridgemse <- min(ridgecv$cvm)
ridgemse <- mean((predict(ridgecv, as.matrix(X), s = "lambda.min") - y)^2)

ridgemse

## [1] 0.5171794

#lassomse <- min(lassocv$cvm)
lassomse <- mean((predict(lassocv, as.matrix(X), s = "lambda.min") - y)^2)

lassomse

## [1] 0.5304655

#pcrmse <- mean(pcr_model$residuals^2)
pcrmse <- mean((predict(pcr_model, as.matrix(X), s = "lambda.min") - y)^2)
pcrmse

## [1] 0.5735587

#plsmse <- mean(plsr_model$residuals^2)
plsmse <- mean((predict(plsr_model, as.matrix(X), s = "lambda.min") - y)^2)
plsmse

## [1] 0.5416171

Test_Error <- c(LSmse, subsetmse, ridgemse, lassomse, pcrmse, plsmse )

Test_Error

## [1] 0.5491941 0.5483947 0.5171794 0.5304655 0.5735587 0.5416171

table <- rbind(table, Test_Error)

table

```



```
##           LS Best_Subset      Ridge      Lasso      PCR
## (Intercept) 2.45234509 2.4523451 2.452345085 2.45234509 2.45234509
## lcavol      0.71640701 0.7798589 0.317607465 0.56536414 0.71640701
## lweight     0.29264240 0.3519101 0.207749630 0.19936439 0.29264240
## age        -0.14254963 0.0000000 -0.003245594 0.00000000 -0.14254963
## lbph       0.21200760 0.0000000 0.132737031 0.02917609 0.21200760
## svi        0.30961953 0.0000000 0.192527220 0.11532140 0.30961953
## lcp       -0.28900562 0.0000000 0.067982376 0.00000000 -0.28900562
## gleason    -0.02091352 0.0000000 0.052456497 0.00000000 -0.02091352
## pgg45      0.27734595 0.0000000 0.109563765 0.01647872 0.27734595
## Test_Error 0.54919414 0.5483947 0.517179403 0.53046551 0.57355874
##           PLS
## (Intercept) 2.45234509
## lcavol      0.71640701
## lweight     0.29264240
## age        -0.14254963
## lbph       0.21200760
## svi        0.30961953
## lcp       -0.28900562
## gleason    -0.02091352
## pgg45      0.27734595
## Test_Error 0.54161713
```

```
cat("Best model is ", colnames(table)[which.min(Test_Error)])
```

```
## Best model is Ridge
```