

Lab 6: Regression with Dimension Reduction

Methods PCR and PLSR

Prof. Gaston Sanchez

Stat 154, Fall 2017

```
read_chunk('lab06-pcr-pls-regression-chunks.R')
```

Introduction

In this lab, you are going to write R code to implement Principal Component Regression (PCR), as well as Partial Least Squares Regression (PLSR). You will also be using the data `Hitters` from the package "ISLR". More specifically, you will use `Salary` as the response variable, and the rest of the variables in `Hitters` as the predictors.

Data Hitters

The data set `Hitters` is part of the R package "ISLR".

```
data(Hitters)
```

```
str(Hitters, vec.len = 1)
```

```
## 'data.frame':   322 obs. of  20 variables:
## $ AtBat      : int  293 315 ...
## $ Hits       : int  66 81 ...
## $ HmRun      : int   1 7 ...
## $ Runs       : int  30 24 ...
## $ RBI        : int  29 38 ...
## $ Walks      : int  14 39 ...
## $ Years      : int   1 14 ...
## $ CAtBat     : int 293 3449 ...
## $ CHits      : int  66 835 ...
## $ CHmRun     : int   1 69 ...
## $ CRuns      : int  30 321 ...
## $ CRBI       : int  29 414 ...
## $ CWalks     : int  14 375 ...
## $ League     : Factor w/ 2 levels "A","N": 1 2 ...
## $ Division   : Factor w/ 2 levels "E","W": 1 2 ...
## $ PutOuts    : int 446 632 ...
## $ Assists    : int  33 43 ...
```

```
## $ Errors      : int  20 10 ...
## $ Salary      : num  NA 475 ...
## $ NewLeague: Factor w/ 2 levels "A","N": 1 2 ...
```

Principal Components Regression (PCR)

Principal Components Regression can be performed with the function `pcr()` which is part of the package "pls". The code below computes PCR for the regression of `Salary` on the rest of 19 predictors.

```
# principal component regression
pcr_fit <- pcr(Salary ~ ., data = Hitters, scale = TRUE, validation = "none")
names(pcr_fit)
```

```
## [1] "coefficients" "scores"      "loadings"    "Yloadings"
## [5] "projection"   "Xmeans"      "Ymeans"      "fitted.values"
## [9] "residuals"    "Xvar"        "Xtotvar"     "fit.time"
## [13] "na.action"    "ncomp"       "method"      "scale"
## [17] "call"         "terms"       "model"
```

1) Start with PCA

You are going write R code in order to replicate the results of `pcr()`. Follow the list of steps shown below:

- Remove observations from `Hitters` that have missing values in `Salary`

```
# remove missing values
hitters <- na.omit(Hitters)
```

- Use `model.matrix()` to create a design matrix based on the formula "`Salary ~ .`".

```
# model matrix
MM <- model.matrix(Salary ~ ., data = hitters)
```

- Note that the generated model matrix includes a constant column for the intercept term. Do not use this column.
- The model matrix (without constant column) will be the matrix of responses. Standardize the model matrix of responses; this will be **X**

```
# exclude 1st column of model matrix
X <- scale(MM[, -1])
```

- The variable `Salary` will be the response **y**

```
y <- hitters$Salary
```

- Use `svd()` to get the Singular Value Decomposition of $\mathbf{X} = \mathbf{U}\mathbf{D}\mathbf{V}^T$

```
# SVD on X (without intercept term)
SVD <- svd(X)
U <- SVD$u
V <- SVD$v
D <- diag(SVD$d)
```

- Compute principal components \mathbf{Z} from the standardized model matrix \mathbf{X} and the eigenvectors in \mathbf{V}

$$\mathbf{Z} = \mathbf{X}\mathbf{V}$$

```
# PCs
Z <- U %*% D
```

- Confirm that your principal components match those of `pcr_fit$scores`

```
# compare Z vs pcr-scores
head(cbind(pcr_fit$scores[,1], Z[,1]))
```

```
##                [,1]        [,2]
## -Alan Ashby      -0.009630358 -0.009630358
## -Alvin Davis      0.410650757  0.410650757
## -Andre Dawson     3.460224766  3.460224766
## -Andres Galarraga -2.553449083 -2.553449083
## -Alfredo Griffin  1.025746581  1.025746581
## -Al Newman       -3.973081710 -3.973081710
```

```
head(cbind(pcr_fit$scores[,19], Z[,19]))
```

```
##                [,1]        [,2]
## -Alan Ashby      -0.019213448 -0.019213448
## -Alvin Davis     -0.003388059 -0.003388059
## -Andre Dawson    -0.066374190 -0.066374190
## -Andres Galarraga -0.013555820 -0.013555820
## -Alfredo Griffin -0.070900042 -0.070900042
## -Al Newman       0.025016886  0.025016886
```

2) PC Regression on the first component

- Use the first PC \mathbf{z}_1 to compute the regression of \mathbf{y} on \mathbf{z}_1 . That is, obtain the first PCR coefficient b_1 given by:

$$b_1 = (\mathbf{z}_1^T \mathbf{z}_1)^{-1} \mathbf{z}_1^T \mathbf{y}$$

```
# regression with first component
pcreg_1 <- lm(y ~ Z[,1])
b1_pcr <- pcreg_1$coefficients
```

- Compute the vector of predicted values $\hat{\mathbf{y}}$:

$$\hat{\mathbf{y}} = b_1 \mathbf{z}_1$$

- Compare your computed $\hat{\mathbf{y}}$ against `pcr_fit$fitted.values[, ,1]`, which is the fitted response using PC1 provided by `pcr()`. Add the average of y to your predicted value before comparison.

```
# compare y-hat with pcr() output
head(cbind(pcreg_1$fitted.values, pcr_fit$fitted.values[, ,1]))

##      [,1]      [,2]
## 1 534.8996 534.8996
## 2 579.6895 579.6895
## 3 904.6869 904.6869
## 4 263.8013 263.8013
## 5 645.2411 645.2411
## 6 112.5090 112.5090
```

3) PC Regression on all PCs

- Compute the vector of PCR-coefficients \mathbf{b}_{pcr} by regressing \mathbf{y} on all principal components \mathbf{Z} :

$$\mathbf{b}_{pcr} = (\mathbf{Z}^T \mathbf{Z})^{-1} \mathbf{Z}^T \mathbf{y}$$

- Compute the vector of predicted values $\hat{\mathbf{y}}$ using all PCs:

$$\begin{aligned} \hat{\mathbf{y}} &= \mathbf{Z}(\mathbf{Z}^T \mathbf{Z})^{-1} \mathbf{Z}^T \mathbf{y} \\ \hat{\mathbf{y}} &= \mathbf{Z} \mathbf{b}_{pcr} \end{aligned}$$

- Compare your computed $\hat{\mathbf{y}}$ against `pcr_fit$fitted.values[, ,19]` and confirm that you have the same results as `pcr()`. Add the average of y to your predicted value before comparison.

```
# compare your coeffs vs those provided by pcr()
head(cbind(pcreg_all$fitted.values, pcr_fit$fitted.values[, ,19]))
```

```
##           [,1]      [,2]
## 1  362.1361  362.1361
## 2  712.6952  712.6952
## 3 1171.3111 1171.3111
## 4   556.7875  556.7875
## 5   493.2515  493.2515
## 6   247.3852  247.3852
```

4) PCR coefficients in terms of the predictor variables

`pcr()` returns regression coefficients—in terms of the predictors—for all possible regressions: with one PC, two PCs, three PCs, and so on, until the regression that uses all 19 PCs.

Consider the PC regression on the first PC \mathbf{z}_1 . The PCR-coefficient is:

$$b_1 = (\mathbf{z}_1^\top \mathbf{z}_1)^{-1} \mathbf{z}_1^\top \mathbf{y}$$

and the fitted $\hat{\mathbf{y}}$ is:

$$\hat{\mathbf{y}} = b_1 \mathbf{z}_1$$

You can re-write the regression of PC1 in terms of the response variables as:

$$\begin{aligned} \hat{\mathbf{y}} &= b_1 \mathbf{z}_1 \\ &= b_1 \mathbf{X} \mathbf{v}_1 \\ &= \mathbf{X} (b_1 \mathbf{v}_1) \\ &= \mathbf{X} \mathbf{b}_1^* \end{aligned}$$

where:

- \mathbf{v}_1 is the loading associated to the first PC, that is, the first column of \mathbf{V}
- \mathbf{b}_1^* is a vector of regression coefficients in terms of the predictors

In general, the PC regression coefficients can be expressed in terms of the predictors as:

$$\mathbf{b}_k^* = \mathbf{V}_k \mathbf{D}_k^{-1} \mathbf{U}_k^\top \mathbf{y}$$

where the index k indicates matrices associated to the first k components. More specifically, \mathbf{V}_k is a matrix of the first k columns of \mathbf{V} , \mathbf{U}_k is a matrix of the first k columns of \mathbf{U} , and \mathbf{D}_k is a $k \times k$ diagonal matrix.

Your turn:

- Take your previously computed coefficient b_1 and calculate the associated vector of coefficients $\mathbf{b}_1^* = b_1 \mathbf{v}_1$. Confirm that your vector \mathbf{b}_1^* matches that of `pcr_fit$coefficients[, 1]`

```
# PCR coeffs with PC1, in terms of predictors
pcreg_1$coefficients[-1] * V[,1]
```

```
## [1] 21.13207878 20.87321071 21.77988064 21.13705999 25.06279956
## [6] 22.26529508 30.11445915 35.21789413 35.24760132 33.99408860
## [11] 36.04328244 36.27081015 33.76212997 -5.80503669 -2.74157997
## [16] 8.28029613 -0.08969488 -0.83758395 -4.46643991
```

```
# compare with output from pcr()
pcr_fit$coefficients[,1]
```

```
##      AtBat      Hits      HmRun      Runs      RBI      Walks
## 21.13207878 20.87321071 21.77988064 21.13705999 25.06279956 22.26529508
##      Years    CAtBat    CHits    CHmRun    CRuns    CRBI
## 30.11445915 35.21789413 35.24760132 33.99408860 36.04328244 36.27081015
##      CWalks    LeagueN    DivisionW    PutOuts    Assists    Errors
## 33.76212997 -5.80503669 -2.74157997 8.28029613 -0.08969488 -0.83758395
## NewLeagueN
## -4.46643991
```

- Do the same for all possible sets of PCs, and verify your coefficients against the output of `pcr_fit$coefficients`.

```
coeffs <- matrix(0, nrow = ncol(X), ncol = ncol(Z))
for (k in 1:ncol(Z)) {
  pcreg_k <- lm(y ~ Z[,1:k])
  if (k == 1) {
    coeffs[,k] <- pcreg_k$coefficients[-1] * V[,1]
  } else {
    coeffs[,k] <- V[,1:k] %*% pcreg_k$coefficients[-1]
  }
}
```

```
rownames(coeffs) <- colnames(X)
colnames(coeffs) <- paste0('Z', 1:ncol(Z))
```

```
# compare against pcr()
cbind(coeffs[,1:3], pcr_fit$coefficients[,1:3])
```

```
##              Z1              Z2              Z3      1 comps      2 comps      3 comps
## AtBat      21.13207878 29.438966 31.596172 21.13207878 29.438966 31.596172
## Hits       20.87321071 29.039128 30.841116 20.87321071 29.039128 30.841116
```

## HmRun	21.77988064	26.912608	21.650526	21.77988064	26.912608	21.650526
## Runs	21.13705999	29.312723	28.894882	21.13705999	29.312723	28.894882
## RBI	25.06279956	31.870731	30.091792	25.06279956	31.870731	30.091792
## Walks	22.26529508	27.235049	28.345853	22.26529508	27.235049	28.345853
## Years	30.11445915	24.434849	25.276570	30.11445915	24.434849	25.276570
## CAtBat	35.21789413	31.042550	33.076799	35.21789413	31.042550	33.076799
## CHits	35.24760132	31.288812	33.388213	35.24760132	31.288812	33.388213
## CHmRun	33.99408860	31.260422	29.160504	33.99408860	31.260422	29.160504
## CRuns	36.04328244	32.314419	33.604363	36.04328244	32.314419	33.604363
## CRBI	36.27081015	32.632509	32.997441	36.27081015	32.632509	32.997441
## CWalks	33.76212997	29.599532	30.624769	33.76212997	29.599532	30.624769
## LeagueN	-5.80503669	-7.865898	5.466047	-5.80503669	-7.865898	5.466047
## DivisionW	-2.74157997	-3.535498	-3.929845	-2.74157997	-3.535498	-3.929845
## PutOuts	8.28029613	11.651168	12.899211	8.28029613	11.651168	12.899211
## Assists	-0.08969488	3.560723	13.245133	-0.08969488	3.560723	13.245133
## Errors	-0.83758395	3.507803	12.826601	-0.83758395	3.507803	12.826601
## NewLeagueN	-4.46643991	-6.145712	7.109718	-4.46643991	-6.145712	7.109718

The lab continues on the next page.

Partial Least Squares Regression

Below are the steps of the PLSR algorithm (in its “classic” version). Assume that the predictors in \mathbf{X} and the response \mathbf{y} are standardized: mean = 0, variance 1.

```
Set  $\mathbf{X}_0 = \mathbf{X}$  and  $\mathbf{y}_0 = \mathbf{y}$ 
for  $h = 1, 2, \dots, r$  do
   $\mathbf{w}_h = \mathbf{X}_{h-1}^\top \mathbf{y}_{h-1}$ 
  normalize weights:  $\|\mathbf{w}_h\| = 1$ 
   $\mathbf{z}_h = \mathbf{X}_{h-1} \mathbf{w}_h / \mathbf{w}_h^\top \mathbf{w}_h$ 
   $\mathbf{p}_h = \mathbf{X}_{h-1}^\top \mathbf{z}_h / \mathbf{z}_h^\top \mathbf{z}_h$ 
   $\mathbf{X}_h = \mathbf{X}_{h-1} - \mathbf{z}_h \mathbf{p}_h^\top$ 
   $b_h = \mathbf{y}_{h-1}^\top \mathbf{z}_h / \mathbf{z}_h^\top \mathbf{z}_h$ 
   $\mathbf{y}_h = \mathbf{y}_{h-1} - b_h \mathbf{z}_h$ 
end for
```

where r is the rank of \mathbf{X}

Your mission is to write R code that carries out PLS regression according to the steps shown above. Your code should contain the following objects:

- **components**: matrix of PLS components \mathbf{Z}
- **weights**: matrix of weights \mathbf{W}
- **loadings**: matrix of loadings \mathbf{P}
- **coefficients**: vector of regression coefficients \mathbf{b}
- **fitted**: matrix of fitted (predicted) values $\hat{\mathbf{Y}}$

The first steps are the same as with PCR:

- Remove observations from `Hitters` that have missing values in `Salary`
- Use `model.matrix()` to create a design matrix based on the formula "`Salary ~ .`".
- Note that the generated model matrix includes a constant column for the intercept term. Do not use this column.
- The model matrix (without constant column) will be the matrix of responses.
- Standardize the model matrix of responses; this will be \mathbf{X}
- The response `Salary` will be \mathbf{y}

Check your first PLS component

- Calculate \mathbf{w}_1 , \mathbf{z}_1 , and \mathbf{p}_1
- Compare your results with `pls_fit$loading.weights[,1]`, `pls_fit$scores[,1]`, `pls_fit$loadings[,1]`,
- Compare the first fitted $\hat{\mathbf{y}}$, i.e. regressing \mathbf{y} on the first PLS component \mathbf{z}_1 , and compare it with `pls_fit$fitted.values[,1]`


```

pls_fit <- pls(Salary ~ ., data = Hitters, scale = TRUE, validation = "none")

pls_regression <- function(X, y) {
  # Assume X is standardized.

  n <- nrow(X)
  p <- ncol(X)
  r <- qr(X)$rank
  Z <- matrix(0, nrow=n, ncol=r)
  W <- matrix(NA, nrow=p, ncol=r)
  P <- matrix(NA, nrow=p, ncol=r)
  b <- rep(0, r)
  Yhat <- matrix(mean(y), nrow=n, ncol=r)

  for (h in 1:r) {
    W[, h] <- t(X) %*% y
    W[, h] <- W[, h] / sqrt(sum(W[, h]^2))
    Z[, h] <- X %*% W[, h]
    P[, h] <- t(X) %*% Z[, h] / sum(Z[, h]^2)
    X <- X - Z[, h] %*% t(P[, h])
    b[h] <- sum(y * Z[, h]) / sum(Z[, h]^2)
    y <- y - b[h] * Z[, h]
    Yhat[, h] <- Yhat[, h] + Z[, 1:h, drop=FALSE] %*% b[1:h]
  }

  list(components=Z, weights=W, loadings=P,
        coefficients=b, fitted=Yhat)
}

my_pls <- pls_regression(X, y)

# Check if the loading weights match.
sapply(1:qr(X)$rank, function(j) {
  all(abs(my_pls$weights[, j] - pls_fit$loading.weights[, j]) < 1e-6)
})

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [15] TRUE TRUE TRUE TRUE TRUE

head(cbind(my_pls$weights[, 1], pls_fit$loading.weights[, 1]))

##           [,1]      [,2]
## AtBat 0.2256137 0.2256137
## Hits  0.2507049 0.2507049
## HmRun 0.1960424 0.1960424

```

```

## Runs 0.2399514 0.2399514
## RBI 0.2568671 0.2568671
## Walks 0.2536725 0.2536725

# Check if the scores match.
sapply(1:qr(X)$rank, function(j) {
  all(abs(my_pls$components[, j] - pls_fit$scores[, j]) < 1e-6)
})

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [15] TRUE TRUE TRUE TRUE TRUE

head(cbind(my_pls$components[, 1], pls_fit$scores[, 1]))

##           [,1]      [,2]
## -Alan Ashby -0.1090169 -0.1090169
## -Alvin Davis 0.6670947 0.6670947
## -Andre Dawson 3.4717021 3.4717021
## -Andres Galarra -2.1298594 -2.1298594
## -Alfredo Griffin 0.9770842 0.9770842
## -Al Newman -4.0036686 -4.0036686

# Check if the loadings match.
sapply(1:qr(X)$rank, function(j) {
  all(abs(my_pls$loadings[, j] - pls_fit$loadings[, j]) < 1e-6)
})

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [15] TRUE TRUE TRUE TRUE TRUE

head(cbind(my_pls$loadings[, 1], pls_fit$loadings[, 1]))

##           [,1]      [,2]
## AtBat 0.2256185 0.2256185
## Hits 0.2231972 0.2231972
## HmRun 0.2179161 0.2179161
## Runs 0.2249696 0.2249696
## RBI 0.2566359 0.2566359
## Walks 0.2292001 0.2292001

# Check if the fitted values match.
sapply(1:qr(X)$rank, function(j) {
  all(abs(my_pls$fitted[, j] - as.vector(pls_fit$fitted.values[, , j])) < 1e-6)
})

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [15] TRUE TRUE TRUE TRUE TRUE

```

```
head(cbind(my_pls$fitted[, 1], as.vector(pls_fit$fitted.values[, , 1])))
```

```
##           [,1]      [,2]  
## [1,] 523.82552 523.82552  
## [2,] 609.97025 609.97025  
## [3,] 921.26845 921.26845  
## [4,] 299.52153 299.52153  
## [5,] 644.37762 644.37762  
## [6,]  91.53754  91.53754
```