

Eigenvalue Decomposition

Predictive Modeling & Statistical Learning

Gaston Sanchez

CC BY-SA 4.0

Introduction

Two Special Decompositions

Last time we talked about the Singular Value Decomposition (SVD).

In these slides, we'll talk about a closely related decomposition of SVD: the so-called **Eigenvalue Decomposition (EVD)** or Spectral Decomposition

Recap

Matrix decompositions, also known as matrix factorizations

$$\mathbf{M} = \mathbf{AB} \quad \text{or} \quad \mathbf{M} = \mathbf{ABC}$$

are a means of expressing a matrix as a product of usually two or three **simpler** matrices.

Types of matrices

Two types of matrices

We said that in data analysis we typically concentrate on two types of matrices:

- ▶ general **rectangular** matrices used to represent data tables.
- ▶ **positive semi-definite** matrices used to represent covariance matrices, correlation matrices, and any matrix that results from a crossproduct.

EVD

Eigenvalue Decomposition

- ▶ In contrast with SVD, the eigenvalue decomposition does NOT apply to any matrix.
- ▶ EVD applies to square matrices in general.
- ▶ A special type of square matrices are **symmetric** matrices.
- ▶ In data analysis methods, these matrices usually appear in the form of cross-product association matrices $\mathbf{X}^T\mathbf{X}$ and $\mathbf{X}\mathbf{X}^T$

EVD

Eigenvalue Decomposition

This decomposition applies to **symmetric** matrices such as the cross-product association matrices $\mathbf{X}^T\mathbf{X}$ and $\mathbf{X}\mathbf{X}^T$

Symmetric Matrices

The attractive thing about EVD is that when applied to symmetric matrices the results have a “simple” nice structure.

Eigen-Value Decomposition

EVD

An $n \times n$ **symmetric matrix** \mathbf{M} can be decomposed as:

$$\mathbf{M} = \mathbf{A}\mathbf{B}\mathbf{A}^T$$

where

- ▶ \mathbf{A} is a $n \times p$ column **orthonormal** matrix containing the **eigen-vectors** of \mathbf{M}
- ▶ \mathbf{B} is a $p \times p$ **diagonal** matrix containing the eigen-values of \mathbf{M}

Eigen-Value Decomposition

EVD

A more common notation for EVD is:

$$\mathbf{M} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

where

- ▶ \mathbf{U} is a $n \times p$ column **orthonormal** matrix containing the eigen-vectors of \mathbf{M}
- ▶ $\mathbf{\Lambda}$ is a $p \times p$ **diagonal** matrix containing the eigen-values of \mathbf{M}

EVD

$$\mathbf{M} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

$$\mathbf{M} = \begin{bmatrix} u_{11} & \cdots & u_{1p} \\ u_{21} & \cdots & u_{2p} \\ \vdots & \ddots & \vdots \\ u_{n1} & \cdots & u_{np} \end{bmatrix} \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_p \end{bmatrix} \begin{bmatrix} u_{11} & \cdots & u_{n1} \\ u_{12} & \cdots & u_{n2} \\ \vdots & \ddots & \vdots \\ u_{1p} & \cdots & u_{np} \end{bmatrix}$$

Eigenvectors

Vectors, which under a given transformation \mathbf{M} map into themselves or multiples of themselves, are called **invariant** vectors under that transformation. It follows that such vectors satisfy the relation:

$$\mathbf{M}\mathbf{x} = \lambda\mathbf{x}$$

where λ is a scalar.

Eigenvectors

The matrix equation:

$$\mathbf{M}\mathbf{x} = \lambda\mathbf{x}$$

can be rearranged as follows:

$$\mathbf{M}\mathbf{x} - \lambda\mathbf{x} = \mathbf{0}$$

Eigenvectors

Given

$$\mathbf{M}\mathbf{x} - \lambda\mathbf{x} = \mathbf{0}$$

We can factor out \mathbf{x}

$$(\mathbf{M} - \lambda\mathbf{I})\mathbf{x} = \mathbf{0}$$

Eigenvectors

Obtaining the eigenstructure of a (square) matrix involves solving the **characteristic equation**

$$\det(\mathbf{M} - \lambda_i \mathbf{I}) = 0$$

If \mathbf{M} is of order $n \times n$, then we can obtain n roots of the equation. These roots are called the **eigenvalues**.

EVD in R

eigen() in R

eigen() function

R provides the function `eigen()` to perform an eigenvalue decomposition of a square matrix.

eigen() output

A list with the following components

- ▶ `values` a vector containing the eigenvalues
- ▶ `vectors` a matrix whose columns contain the eigenvectors

EVD example in R

```
# X'X matrix
set.seed(22)
X <- as.matrix(USArrests)
XtX <- t(X) %*% X

# eigenvalue decomposition
EVD = eigen(XtX)

# elements returned by eigen()
names(EVD)

## [1] "values" "vectors"

# vector of eigenvalues
(lambdas = EVD$values)

## [1] 2013735.2431 37957.1103 2084.9578 326.5089
```

EVD example in R (con't)

```
# matrix of eigenvectors
```

```
(V <- EVD$vectors)
```

```
##           [,1]      [,2]      [,3]      [,4]  
## [1,] -0.04239181  0.01616262  0.06588426  0.99679535  
## [2,] -0.94395706  0.32068580 -0.06655170 -0.04094568  
## [3,] -0.30842767 -0.93845891 -0.15496743  0.01234261  
## [4,] -0.10963744 -0.12725666  0.98347101 -0.06760284
```

Relation between EVD and SVD

The EVD of the cross-product matrix of columns (or minor product moment) $\mathbf{X}^T\mathbf{X}$ can be expressed as:

$$\mathbf{X}^T\mathbf{X} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^T$$

in terms of the SVD factorization of \mathbf{X} :

$$\mathbf{X}^T\mathbf{X} = \mathbf{V}\mathbf{D}^2\mathbf{V}^T$$

Relation between EVD and SVD

The EVD of the cross-product matrix of rows (or major product moment) \mathbf{XX}^T can be expressed as:

$$\mathbf{XX}^T = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

in terms of the SVD factorization of \mathbf{X} :

$$\mathbf{XX}^T = \mathbf{U}\mathbf{D}^2\mathbf{U}^T$$

Power Method

About the Power Method

One of the basic procedures following a successive approximation approach is precisely the **Power Method**.

In its simplest form, the Power Method (PM) allows us to find **the largest eigenvector** and its corresponding **eigenvalue**.

About the Power Method

Choose an arbitrary vector \mathbf{w}_0 to which we will apply the symmetric matrix \mathbf{S} repeatedly to form the following sequence:

$$\mathbf{w}_1 = \mathbf{S}\mathbf{w}_0$$

$$\mathbf{w}_2 = \mathbf{S}\mathbf{w}_1 = \mathbf{S}^2\mathbf{w}_0$$

$$\mathbf{w}_3 = \mathbf{S}\mathbf{w}_2 = \mathbf{S}^3\mathbf{w}_0$$

$$\vdots$$

$$\mathbf{w}_k = \mathbf{S}\mathbf{w}_{k-1} = \mathbf{S}^k\mathbf{w}_0$$


Power Method: Example

Consider a matrix \mathbf{S}

$$\mathbf{S} = \begin{bmatrix} 3 & 1 \\ 1 & 2 \end{bmatrix}$$


and an initial vector \mathbf{w}_0

$$\mathbf{w}_0 = \begin{bmatrix} -0.4 \\ 1.3 \end{bmatrix}$$



About the Power Method

- ▶ In practice, we must **rescale** the obtained vector \mathbf{w}_k at each step.
- ▶ The rescaling will allows us to judge whether the sequence is converging.
- ▶ After some iterations, the vector \mathbf{w}_{k-1} and \mathbf{w}_k will be very similar
- ▶ Assuming a reasonable scaling strategy, the sequence will **usually converge to the dominant eigenvector of \mathbf{S} .**



Dominant Eigenvalue

The obtained vector is the dominant eigenvector. To get the corresponding eigenvalue we calculate the so-called **Rayleigh quotient** given by:

$$\lambda = \frac{\mathbf{w}_k^T \mathbf{S} \mathbf{w}_k}{\mathbf{w}_k^T \mathbf{w}_k}$$

Remarks

Conditions for the power method to be successfully used:

- ▶ The matrix must have a *dominant* eigenvalue.
- ▶ The starting vector \mathbf{w}_0 must be nonzero.
- ▶ We need to scale each of the vectors \mathbf{w}_k otherwise the algorithm will “explode”

PM Pseudocode

Let's consider a more detailed version of the PM algorithm:

1. Start with an arbitrary initial vector \mathbf{w}
2. Obtain product $\tilde{\mathbf{w}} = \mathbf{S}\mathbf{w}$
3. Normalize $\tilde{\mathbf{w}}$

$$\text{e.g. } \mathbf{w} = \frac{\tilde{\mathbf{w}}}{\|\tilde{\mathbf{w}}\|_{p=2}}$$

4. Compare \mathbf{w} with its previous version
5. Repeat steps 2 till 4 until convergence

Why does the PM work?

Assume that the matrix \mathbf{S} has p eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_p$, and that they are ordered in decreasing way $|\lambda_1| > |\lambda_2| \geq \dots \geq |\lambda_p|$.

Note that the first eigenvalue is strictly greater than the second one. This is a very important assumption.

In the same way, we'll assume that the matrix \mathbf{S} has p linearly independent vectors $\mathbf{u}_1, \dots, \mathbf{u}_p$ ordered in such a way that \mathbf{u}_j corresponds to λ_j .

Why does the PM work?

The initial vector \mathbf{w}_0 may be expressed as a linear combination of $\mathbf{u}_1, \dots, \mathbf{u}_p$

$$\mathbf{w}_0 = a_1 \mathbf{u}_1 + \dots + a_p \mathbf{u}_p$$

At every step of the iterative process the vector \mathbf{w}_k is given by:

$$\mathbf{w}_k = a_1 \lambda_1^k \mathbf{u}_1 + \dots + a_p \lambda_p^k \mathbf{u}_p$$

Why does the PM work?

Since λ_1 is the dominant eigenvalue, the component in the direction of \mathbf{u}_1 becomes relatively greater than the other components as k increases. If we knew λ_1 in advance, we could rescale at each step by dividing by it to get:

$$\left(\frac{1}{\lambda_1^k}\right) \mathbf{w}_k = a_1 \mathbf{u}_1 + \cdots + a_p \left(\frac{\lambda_p^k}{\lambda_1^k}\right) \mathbf{u}_p$$

which converges to the eigenvector $a_1 \mathbf{u}_1$, provided that a_1 is nonzero.

Why does the PM work?

Of course, in real life this scaling strategy is not possible—we don't know λ_1 . Consequently, the eigenvector is determined only up to a constant multiple, which is not a concern since the really **important thing is the *direction*** not the length of the vector.

The speed of the convergence depends on how bigger λ_1 is respect with to λ_2 , and on the choice of the initial vector \mathbf{w}_0 . If λ_1 is not much larger than λ_2 , then the convergence will be slow.

More Remarks

- ▶ The power method is a sequential method.
- ▶ We can obtain $\mathbf{w}_1, \mathbf{w}_2$, and so on, step by step.
- ▶ If we only need the first k vectors, we can stop the procedure at the desired stage.

Obtaining more eigenvectors?

Once we've obtained the first eigenvector \mathbf{w}_1 and eigenvalue λ_1 , we can compute the second eigenvector by reducing the matrix \mathbf{S} by the amount explained by the first eigenvector.

This operation of reduction is called **deflation** and the residual matrix is obtained as:

$$\mathbf{S}_1 = \mathbf{S} - \lambda_1 \mathbf{w}_1 \mathbf{w}_1^T$$

To get the second eigenvalue and its corresponding eigenvector, we operate on \mathbf{S}_1 in the same way as the operations on \mathbf{S} .